

CS 240 Homework 2

Lan Liu 7956394

Pritha D N 9727371

Experiment 1: Strong scaling analysis

1. Set $k=864$, as a common multiple of $p=1,4,8,12,16,24,48,72$.

2. Use compute node for each running.

3. Stop condition: $\text{relres} < 1e-8$.

$k=864$ is fairly large, based on hw2harness.h, *If you want to verify correctly on large problems ($k > 400$ or so), then reduce the norm of the residual requirement by an order of magnitude or two.*

4. After 2198 iterations, cgsolve stops and the norm of residual is $3.77e-7$.

Remark: Here we did not use value $p=32$ and $p=64$, instead, we use $p=48$ and 72 , which are multiple of 24 . We have two reasons to do this:

Reason 1: save SUs.

For each compute node, the total processor number is 24 , and no matter how many processors we actually use, the SUs consumed will always be counted as 24 . For $p=64$, it requires 4 compute nodes with 16 processor on each node, and need $4*24=96$ SUs. But for $p=72$, we only need 3 nodes with 24 processors on each node, and need $3*24=72$ SUs. And multiple of 24 will make full use of processor and save SUs. [This was told by Burak]

Reason2: save running time.

From the result below, we can see that $p=64$ takes longer time than $p=48, 24, 16$. It is because communication within a node is faster than communication between nodes, and it is also cheaper. So we should set p be multiples of 24 , i.e. $24, 48$ and 72 .

Number of processors	Time taken	Efficiency
1	24.056	
4	6.979	0.862
8	4.084	0.736
12	2.839	0.706
16	2.118	0.710
24	1.47	0.682
48	0.823	0.609
72	0.680	0.491
32	1.117	0.673
64	2.910	0.129

Form1: Result for $k=864$, iteration steps =2198, residual= $3.77e-7$

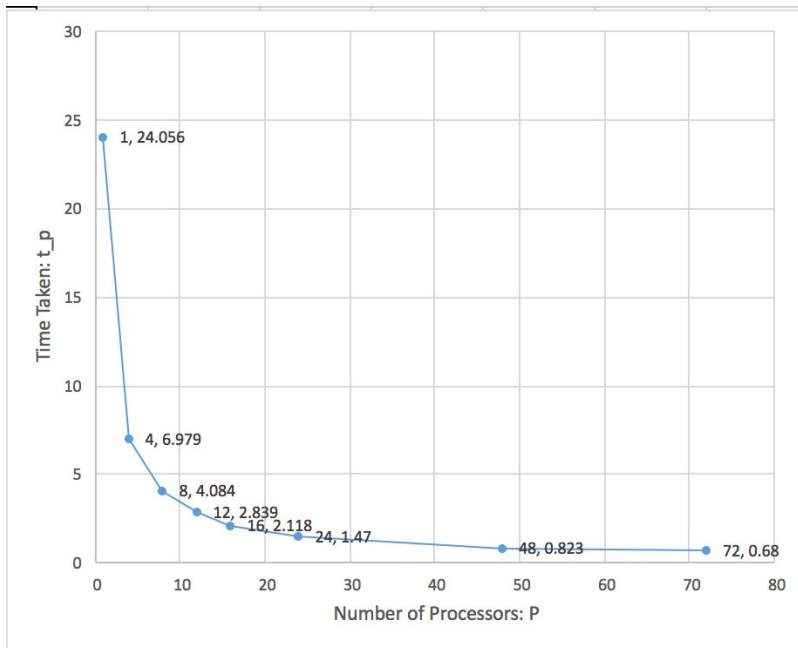


Figure 1: Running time versus number of processors P




Figure 2: Efficiency versus number of processors P

As number of processors increase, the efficiency will decrease, this match with Latency/ Bandwidth Model. The concrete analysis will be on the last part.

Experiment 2: Weak scaling analysis

K	Number of processors	Time taken	Efficiency
864	1	1.088321	
1728	4	1.178593	0.231
2440	8	1.417219	0.096
2988	12	1.539454	0.059
3456	16	1.818838	0.0373
4224	24	3.114038	0.0145
6000	48	3.648195	0.00621
7344	72	Times out	

Form 2: K is the closest number proportional to \sqrt{p} and also multiple of p. Do 100 iterations.




Figure 3: Running time versus number of processors P




Figure 4: Efficiency versus number of processors P

Experiment 3: TAU analysis of the performance

K = 144; p = 8; maxiterations = 10

NODE 0;CONTEXT 0;THREAD 0:

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	0.741	940	1	77	940066 int main(int, char **) C
84.2	791	791	1	0	791795 MPI_Init()
15.4	145	145	1	0	145113 void save_vec(int, double *) C
0.2	1	1	1	0	1544 MPI_Finalize()
0.0	0.35	0.35	30	0	12 MPI_Bcast()
0.0	0.273	0.273	21	0	13 MPI_Reduce()
0.0	0.162	0.162	1	0	162 MPI_Gather()
0.0	0.064	0.064	10	0	6 MPI_Recv()
0.0	0.023	0.023	10	0	2 MPI_Send()
0.0	0.001	0.001	1	0	1 MPI_Comm_size()
0.0	0	0	1	0	0 MPI_Comm_rank()

USER EVENTS Profile :NODE 0, CONTEXT 0, THREAD 0

NumSamples	MaxValue	MinValue	MeanValue	Std. Dev.	Event Name
30	8	8	8	0	Message size for broadcast
1	2.074E+04	2.074E+04	2.074E+04	0	Message size for gather
21	8	8	8	0	Message size for reduce

NODE 2;CONTEXT 0;THREAD 0:

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	0.721	939	1	97	939992 int main(int, char **) C
84.2	791	791	1	0	791775 MPI_Init()
14.9	140	140	1	0	140416 void save_vec(int, double *) C
0.7	6	6	1	0	6350 MPI_Finalize()
0.0	0.458	0.458	30	0	15 MPI_Bcast()
0.0	0.155	0.155	21	0	7 MPI_Reduce()
0.0	0.059	0.059	20	0	3 MPI_Recv()
0.0	0.041	0.041	20	0	2 MPI_Send()
0.0	0.016	0.016	1	0	16 MPI_Gather()
0.0	0.001	0.001	1	0	1 MPI_Comm_size()
0.0	0	0	1	0	0 MPI_Comm_rank()

USER EVENTS Profile :NODE 2, CONTEXT 0, THREAD 0

NumSamples	MaxValue	MinValue	MeanValue	Std. Dev.	Event Name
30	8	8	8	0	Message size for broadcast
0	0	0	0	0	Message size for gather
21	8	8	8	0	Message size for reduce


FUNCTION SUMMARY (total):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	5	7,520	8	736	940029 int main(int, char **) C
84.2	6,334	6,334	8	0	791763 MPI_Init()
15.3	1,149	1,149	8	0	143682 void save_vec(int, double *) C
0.3	24	24	8	0	3092 MPI_Finalize()
0.0	3	3	240	0	15 MPI_Bcast()
0.0	1	1	168	0	7 MPI_Reduce()
0.0	0.727	0.727	140	0	5 MPI_Recv()
0.0	0.345	0.345	140	0	2 MPI_Send()
0.0	0.293	0.293	8	0	37 MPI_Gather()
0.0	0.006	0.006	8	0	1 MPI_Comm_size()
0.0	0.002	0.002	8	0	0 MPI_Comm_rank()

FUNCTION SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	0.738	940	1	92	940029 int main(int, char **) C
84.2	791	791	1	0	791763 MPI_Init()
15.3	143	143	1	0	143682 void save_vec(int, double *) C
0.3	3	3	1	0	3092 MPI_Finalize()
0.0	0.445	0.445	30	0	15 MPI_Bcast()
0.0	0.138	0.138	21	0	7 MPI_Reduce()
0.0	0.0909	0.0909	17.5	0	5 MPI_Recv()
0.0	0.0431	0.0431	17.5	0	2 MPI_Send()
0.0	0.0366	0.0366	1	0	37 MPI_Gather()
0.0	0.00075	0.00075	1	0	1 MPI_Comm_size()
0.0	0.00025	0.00025	1	0	0 MPI_Comm_rank()

Name: MPI_Finalize()
Metric Name: TIME
Value: Exclusive
Units: seconds



Metric: TIME
Value: Exclusive



Blue - MPI_Init; Red - save_vec; Green - MPI_Finalize

Per iteration:

1. Total number of send/receive = $2(p-1)$, since first and last processor only send one set of k elements. Other processors send $2k$ elements.
2. Total broadcasts = $3p$ (alpha, beta, relres)
3. Total reduce = $2p$ [We initialize rtr before the loop begins which counts as p computations, independent of iterations]
4. Total gather = p

4. Analysis with Communication Volume Model

4.1 Sequential Matlab Code

```

while relres > 1e-6 && niters < maxiters
    niters = niters+1;
    Ad = matvec(d,n);      % MATVEC
    alpha = rtr / (d'*Ad); % DDOT
    x = x + alpha * d;      % SAXPY
    r = r - alpha * Ad;      % SAXPY
    rtrold = rtr;
    rtr = r'*r;             % DDOT
    beta = rtr / rtrold;
    d = r + beta * d;        % SAXPY
    relres = sqrt(rtr) / normb;
end;

```

P=1, for each iteration step, the communication cost is 0, and computation cost is (only count product and division, ignore lower order):

1 MATVEC: 5n

3 SAXPY: 3n

2 DDOT: 2n

Total cost per step is 10n.

4.2 Parallel Code

```

while(relres > 1e-8 && *niters < maxiterations)
{
    *niters = *niters+1;
    Ad = matvec(d, n, rank, p);
    dot_product = ddot(d,Ad,n,p);
    MPI_Reduce(&dot_product,&dAd,1,MPI_DOUBLE,MPI_SUM,0,MPI_COMM_WORLD);
    if(rank==0)
    {
        alpha = rtr / dAd;
    }
    MPI_Bcast(&alpha,1, MPI_DOUBLE,0,MPI_COMM_WORLD);
    x = saxpy(alpha,x,d,n,p);
    r = saxpy(-1*alpha,r,Ad,n,p);
    if(rank==0)
    {
        rtrold = rtr;
    }
    dot_product=ddot(r,r,n,p);
    MPI_Reduce(&dot_product,&rtr,1,MPI_DOUBLE,MPI_SUM,0,MPI_COMM_WORLD);

    if(rank==0)
    {
        beta = rtr / rtrold;
    }
    MPI_Bcast(&beta,1, MPI_DOUBLE,0,MPI_COMM_WORLD);
    d = saxpy(beta,r,d,n,p);
    if(rank==0)
    {
        relres = sqrt(rtr) / normb;
    }
    MPI_Bcast(&relres,1, MPI_DOUBLE,0,MPI_COMM_WORLD);
}

```

	computation cost		communication cost
1 MATVEC	$5n/p$	$\sim 2p$ MPI_Send/Receive	$2k*(p-2)+k*2$
3 SAXPY	$3n/p$	3 MPI_Bcast	$3(p-1)$
2 DDOT	$2n/p$	2 MPI_Reduce	$2(p-1)$
Total	$10n/p$	1 MPI_Gather	n
		Total	$2kp+n$

Note: For Matvec function, each processor other than first and last should send twice k elements and receive twice k elements, and the first and last processor should send and receive once k elements. So the total cost is $2k*(p-2)+k*2$

As k increase, the computation cost will decrease and the communication cost will increase. For experiment 2, if we set k be about proportional to \sqrt{p} , then computation cost $n/p=k^2/p$ is about a constant, but the communication cost will increase as k and p increase, so we will see a pattern that as p and k increase, the running time is increasing.